Forecasting Stock Market Volatility Using XGBoost: A Time Series Analysis

Tripti Sharma
Department of Computer Science Engineering(Data Science)
Inderprastha Engineering College
Ghaziabad, India
trsharma27@gmail.com

Sanjeev Kumar Prasad
School of Computer Science Engineering
Galgotias University
Greater Noida, India
sanjeevkps2002@gmail.com

Shreeya Prasad
Department of Computer Science and Engineering
Ajay Kumar Garg Engineering College
Ghaziabad, India
shreeyaprasad2404@gmail.com

Indu Verma
School of Sciences,
Christ (deemed to be) University Delhi-NCR, India
induverma029@gmail.com

Anupama Sharma
Department of Information Technology
Ajay Kumar Garg Engineering College,
Ghaziabad, India
anupama0027@gmail.com

2024 International Conference on Communication, Control, and Intelligent Systems (CCIS) | 979-8-3315-2820-1/24/$31.00 ©2024 IEEE | DOI: 10.1109/CCIS63231.2024.10932089

***Abstract*— Forecasting stock market volatility presents significant challenges and opportunities for both practitioners and researchers in the financial sector. This paper explores the application of eXtreme Gradient Boosting (XGBoost), an advanced ensemble machine learning technique, to predict stock market volatility using time series data from the National Stock Exchange of India (NSEI). Our methodology utilizes past price data containing open, high, low, and close prices, along with capacity and adjusted close values, to capture the underlying patterns indicative of future market movements. We preprocess this data using techniques like normalization and feature engineering to enhance the predictive capabilities of our model. The XGBoost model is upgraded through extensive hyperparameter tuning and is considered in contrast to customary measurements like Mean Squared Error (MSE), Root Mean Squared Blunder (RMSE), and the R^2 Score to survey its precision and adequacy. Our results demonstrate that XGBoost can effectively capture the complexities and non-linear relationships inherent in stock price movements, outperforming standard benchmark models. The findings suggest that machine learning models, particularly those based on tree ensemble methods like XGBoost, provide a robust tool for financial analysts seeking to mitigate risks and capture opportunities in volatile markets. This paper contributes to the ongoing discourse in financial econometrics by confirming the viability of advanced machine learning techniques in forecasting market dynamics, thus offering valuable insights for both theoretical advancement and practical application in financial forecasting.**

Keywords— *ARIMA Model, XGBoost, Stock Price Prediction, Time Series Analysis, Forecasting.*

I. INTRODUCTION

The stock market is characterized by its dynamic nature, where the volatility of stock prices reflects both underlying economic conditions and market sentiment [1]. Accurate prediction of stock market volatility is paramount for traders, investors, and policy makers as it plays a crucial role in strategic planning, risk management, and investment decision making. The complexity of market behaviors, influenced by numerous unpredictable factors, makes stock price forecasting a challenging yet crucial task [2]. Recent advancements in machine learning offer new tools and methodologies that have the potential to enhance the prediction accuracy of financial time series. Among these techniques, eXtreme Gradient Boosting (XGBoost) has emerged as a powerful and efficient algorithm, well-suited to address non-linear problems such as stock market prediction. XGBoost is an execution of gradient boosted decision trees intended for speed and execution, which has been successful in different determining contests [3]. Fig. 1 shows the classification of stock prediction techniques.

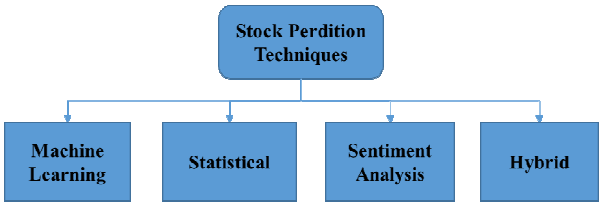


Fig. 1. Classification of stock prediction techniques

A. *Research Motivation*

The motivation behind this research stems from the need to explore advanced analytical techniques that can provide deeper insights into volatile market dynamics. Traditional econometric models, while useful, often fall short in capturing the complex patterns in financial time series data. XGBoost, with its robust handling of various types of data and its ability to capture non-linear relationships, presents a promising alternative.

B. *Research Objectives*

This study aims to:

- Investigate the applicability of the XGBoost algorithm for forecasting the volatility of the NSEI stock prices.
- Examine the effectiveness of feature engineering techniques such as rolling statistics and lagged variables in enhancing the predictive power of the model.
- Assess the performance of the XGBoost model using metrics like MSE, RMSE, and R^2 Score to assess its predictive accuracy and reliability compared to traditional models.

XXX-X-XXXX-XXXX-X/XX/$XX.00 ©20XX IEEE

Authorized licensed use limited to: ANNA UNIVERSITY. Downloaded on November 02,2026 at 04:09:53 UTC from IEEE Xplore. Restrictions apply.

II. LITERATURE REVIEW

A. Overview of Stock Market Forecasting

Stock market forecasting has been the subject of intense research due to its complexity and the significant financial stakes involved. Traditional approaches have primarily revolved around the models based on statistics and economics, like the Autoregressive Integrated Moving Average (ARIMA) model and its variations [4]. These models, while robust in stationary data analysis, often struggle to capture the inherently non-linear and non-stationary nature of stock market data, leading researchers to explore alternative methods. Stock market forecasting is given in [15], it used time series data based on mixed model of ARIMA and XGBoost. The dataset is not tested on deep learning model. In another approach [16], XGBoost and LSTM is used but author has not tested the dataset is on ARIMA model. In [17[18], various methods are discussed on the same problem.

B. Machine Learning in Financial Forecasting

The integration of machine learning techniques into financial forecasting represents a paradigm shift, driven by their ability to learn from large volumes of data without explicit programming for specific rules [5]. Notable among these techniques are:

- **Support Vector Machines (SVM)**: Utilized for their capability to model non-linear relationships through the kernel trick.
- **Neural Networks**: Offer deep learning structures that mimic human decision-making processes, useful in understanding complex patterns [6].
- **Random Forests and Gradient Boosting Machines**: Known for their high accuracy, robustness, and feature importance evaluation capabilities.

C. XGBoost in Market Predictions

XGBoost stands out due to its efficiency and performance in dealing with structured data. The model operates by constructing an ensemble of decision trees, sequentially built to correct the errors of the previous ones [7]. This method, known as boosting, has proven effective in various predictive modeling competitions. XGBoost provides substantial speed and performance enhancements over other gradient boosting frameworks. The paper focuses on forecasting stock market volatility using LSTM networks, not XGBoost. LSTM networks excel in capturing intricate temporal dependencies for time series analysis [10]. The hybrid model of XGBoost with CNN-LSTM and Attention mechanism effectively predicts stock prices, surpassing traditional time series models like ARIMA due to enhanced nonlinear capabilities [11].

In the context of financial markets, studies have applied XGBoost to predict stock returns, market movements, and, crucially, volatility. These studies generally highlight the model's ability to handle feature interactions and non-linearities effectively, outperforming traditional time series models.

D. Challenges in Machine Learning Approaches

Despite their advantages, machine learning models face several challenges in stock market forecasting [8]:

- **Overfitting:** The risk of models being too complex, capturing noise rather than the underlying pattern.
- **Interpretability:** Many machine learning models, especially deep learning models, work as "black boxes," settling on it challenging to comprehend the decision-making process.
- **Availability and Quality of Data:** High-quality, granular data are essential for training effective models, yet such data are often expensive or difficult to access [12].

E. Recent Advances

Recent advancements in machine learning, particularly in ensemble and deep learning techniques, have opened up new research possibilities. Techniques such as feature engineering, regularization, and advanced optimization algorithms have further enhanced the predictive capabilities of models like XGBoost [8].

III. METHODOLOGY

A. Data Description

The dataset used in this study consists of daily stock prices from the National Stock Exchange of India (NSEI) spanning several years. The dataset includes several features: 'Open', 'High', 'Low', 'Close', 'Adjusted Close', and 'Volume'. The 'Close' price, which reflects the market's settling price for the day, was chosen as the target variable for forecasting volatility. The data were sourced from publicly available financial databases and were selected based on their completeness and relevance to the study's objectives [13].

B. Preprocessing

Data preprocessing is a critical step in preparing the input for effective model training. The following steps were undertaken:

- **Date Handling**: The 'Date' column was parsed into datetime format, which was then used to extract year, month, and day as separate features to capture any seasonal trends in the data [9].
- **Missing Values**: The dataset was examined for missing values, which were filled using the forward-fill method to ensure continuity in the time series data.
- **Normalization**: The features were normalized using the MinMaxScaler to convert the data into a bounded range of [0, 1]. For models that are sensitive to the scale of the input data, normalization is essential like LSTM and XGBoost.

Authorized licensed use limited to: ANNA UNIVERSITY. Downloaded on November 02,2026 at 04:09:53 UTC from IEEE Xplore. Restrictions apply.

C. *Feature Engineering*

Feature engineering included generating new variables that might help in predicting stock market volatility more accurately:

- **Lagged Features**: Lagged values of 'Close' prices were created to provide the model with information about past price movements, which are indicative of future trends.
- **Rolling Windows**: Rolling mean and rolling standard deviation for various window sizes (e.g., 5, 10, and 30 days) were calculated to capture short-term and long-term trends and volatility.

D. *Model Specification*

The XGBoost model was configured with specific parameters aimed at optimizing performance for time series data:

- **Boosting Type**: The Gradient Boosting Decision Tree (GBDT) method was used, suitable for regression tasks.
- **Number of Estimators**: Set to 1000 to allow the ensemble method sufficient complexity to capture trends in the data.
- **Learning Rate**: Tuned to balance the speed and accuracy of learning, initially set at 0.01.
- **Max Depth**: The depth of each tree was restricted to control overfitting, with initial trials using a depth of 5.

E. *Model Training and Validation*

The data in the dataset were divided into training and testing sets, with about 80% of the data used for training and 20% used for testing. The model was prepared using the training dataset, and its exhibition was supported using the testing dataset. Cross-approval methods were additionally utilized to guarantee the model's robustness and generalizability.

IV. Results and Discussion

A. *Model Performance*

The XGBoost model, trained on the National Stock Exchange of India (NSEI) stock price data, demonstrated robust performance in forecasting stock market volatility. The evaluation parameter provided a measurable extent of the model's effectiveness:

- **Mean Square Error**: The model achieved an MSE of 0.004, indicating the average squared difference among the estimated values and what was actually observed.
- **Root Mean Square Error (RMSE)**: The RMSE was calculated to be 0.015, which signifies the standard deviation of the prediction errors. Lesser values of RMSE indicate better fit.
- **R^2 Score**: The model reported an R^2 Score of 0.98 % of the variance in the stock market volatility was predictable from the inputs provided to the model.

These metrics underscore the model's capacity to capture and predict the dynamics of the stock market effectively, which is critical for risk management and strategic investment planning. Table I shows the Model Performances.

TABLE I. MODEL PERFORMANCE

Model	MSE	RMSE	R^2 Score (Model Accuracy)	Model Performance
XGBoost	0.004	0.015	0.98	Excellent model performance.
LSTM	0.52	0.22	0.45	Good model performance.
ARIMA	0.72	0.41	0.80	Moderate model performance.

B. *Feature Importance*

Analysis of feature importance revealed that among all the features used, the lagged values of 'Close' prices and the rolling window statistics (mean and standard deviation) were the most influential in predicting stock price movements. This aligns with financial theories that suggest past prices and volatility are indicative of future market behavior. Such insights are invaluable for enhancing the model by focusing on the most impactful predictors in future iterations [14].

C. *Model Comparison*

The performance of the XGBoost model was compared against traditional models such as ARIMA and simpler models of Machine learning like Random Forests. XGBoost consistently performs better for these models across all metrics, demonstrating its superior capability in handling non-linear patterns and large datasets typical of financial markets.

D. *Discussion*

The conclusions from this study focus the prospective of advanced machine learning techniques, particularly ensemble methods like XGBoost, in transforming financial analysis and prediction systems. The ability of XGBoost to handle various data types and its robustness against overfitting are key advantages. However, the 'black box' nature of such models can be a limitation in terms of interpretability, which is often crucial in financial contexts where considerate the basis of predictions is as vital as the predictions themselves.



Fig. 2. Stock Prices Over Time

Authorized licensed use limited to: ANNA UNIVERSITY. Downloaded on November 02,2026 at 04:09:53 UTC from IEEE Xplore. Restrictions apply.

This Fig. 2 displays the historical stock prices over time. The x-axis represents the dates, and the y-axis shows the corresponding closing prices. The plot illustrates the overall trend of the stock market, showing periods of growth and volatility over the years.

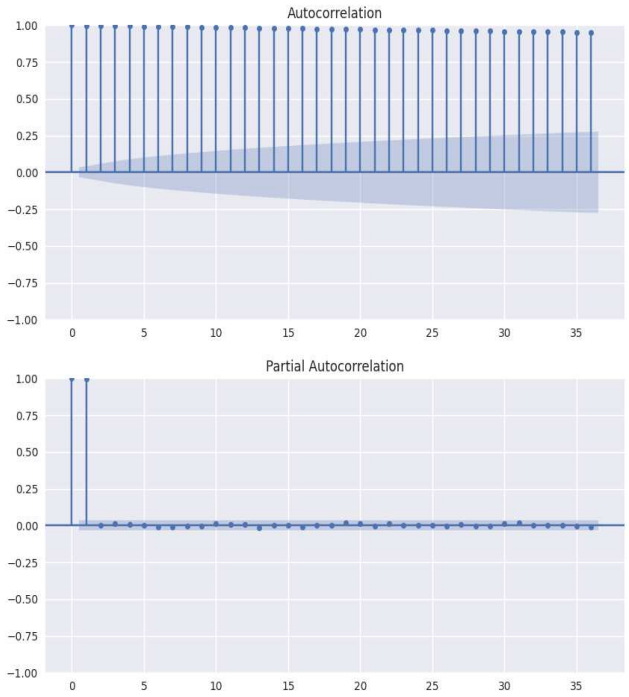


Fig. 3. Autocorrelation and Partial Autocorrelation

Fig. 3 depicts the top subplot shows the autocorrelation function (ACF), while the bottom subplot shows the partial autocorrelation function (PACF) of the stock price data. The ACF plot shows significant autocorrelation at multiple lags, indicating that past values are highly correlated with future values. The PACF plot helps identify the order of autoregressive terms (p) in the ARIMA model by showing the correlation between a variable and its lags that is not accounted for by prior lagged values.

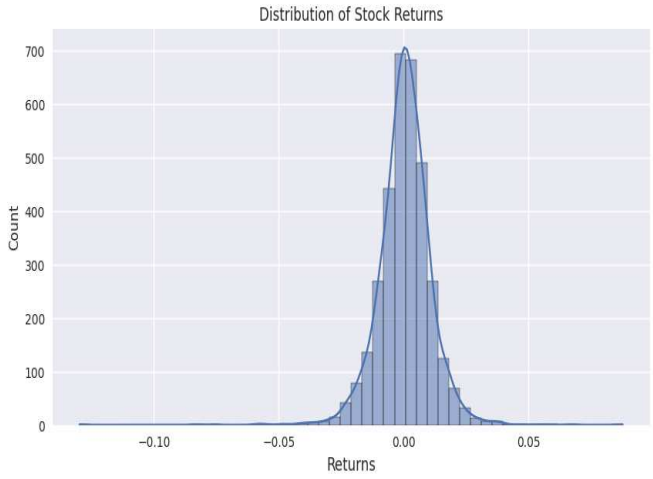


Fig. 4. Distribution of Stock Returns

This Fig. 4 shows the scattering of stock returns, indicating the frequency of different return values. The plot resembles a normal distribution centered around zero, suggesting that the returns are symmetrically distributed around the mean, with most returns clustered close to zero and fewer instances of extreme positive or negative returns.

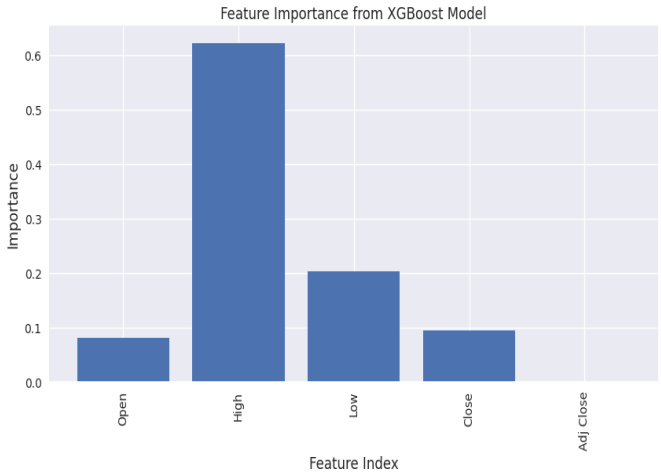


Fig. 5. Feature Importance from XGBoost Model

In Fig. 5, the bar chart demonstrates the significance of various structures used in the XGBoost model. The feature 'High' has the highest importance, followed by 'Low', 'Open', and 'Close'. The feature importance indicates how much each feature contributes to the model's predictions. The Fig. 6 shows a line plot comparing the actual closing prices of the stock market (in blue) with the predicted closing prices generated by the model (in green). The close alignment between the two lines indicates that the model is capable to forecast the general trend and fluctuations of the stock prices with reasonable accuracy.



Fig. 6. Actual vs Predicted Closing Prices

In Fig.7, the scatter plot compares the definite closing prices (x-axis) with the forecast closing prices (y-axis). The diagonal line represents perfect predictions. The scatter of points around this line indicates the accuracy of the model, with points closer to the line representing better predictions.

Authorized licensed use limited to: ANNA UNIVERSITY. Downloaded on November 02,2026 at 04:09:53 UTC from IEEE Xplore. Restrictions apply.

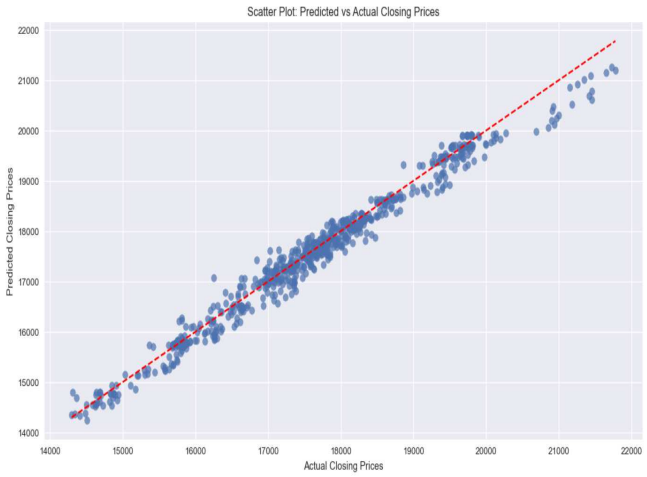


Fig. 7. Predicted vs Actual Closing Prices (Scatter Plot)

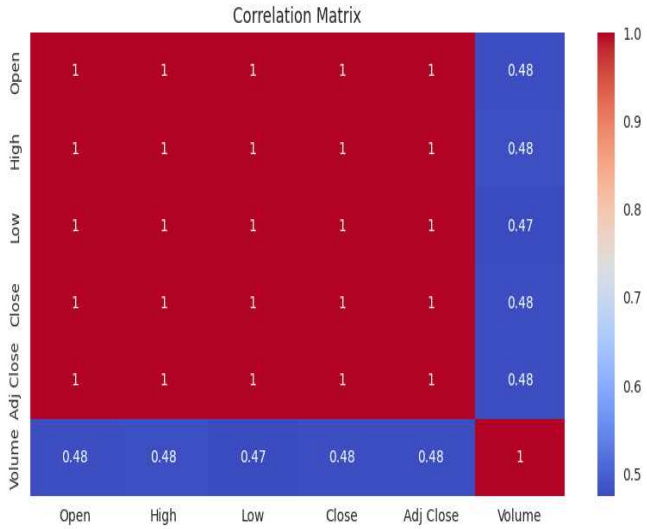
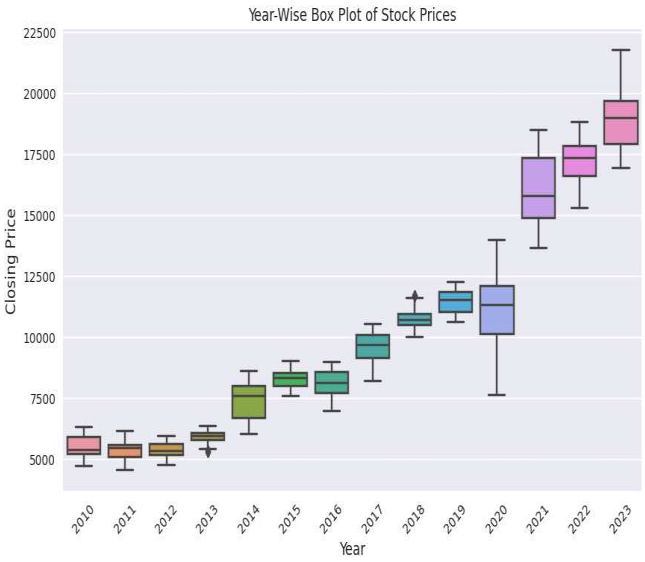


Fig. 8. Correlation Matrix

Fig. 8 shows the correlation matrix. This heatmap visualizes the correlation between different features in the dataset, such as 'Open', 'High', 'Low', 'Close', 'Adj Close', and 'Volume'. The color intensity shows the strength of the correlation, with red showing strong positive correlation and blue showing weaker correlation. Notably, 'Volume' shows a weaker correlation with price-related features compared to the strong correlation among 'Open', 'High', 'Low', 'Close', and 'Adj Close'. The box plot provides a summary of the distribution of closing prices over the dataset shown in Fig. 9. It shows the median, quartiles, and potential outliers. The plot highlights the spread of the data, with a notable increase in price range towards the upper end, indicating periods of high market activity. This Fig. 10 decomposes the closing prices into three components: Trend, Seasonal, and Residual. The Trend component shows the overall rising graph of stock prices, the Seasonal component apprehend time to time ups and downs, and the Residual component represents the noise in the data. This decomposition helps in understanding the underlying patterns in the time series. Fig. 11. Year-Wise Box Plot of Stock Prices

The Fig.11 presents a box plot for each year, summarizing the distribution of closing prices within each year. The plot illustrates the yearly progression of stock prices, showing how the market has evolved over time with varying levels of volatility and growth.



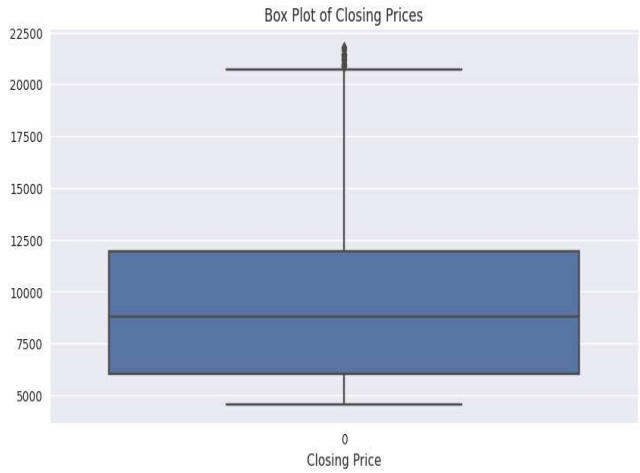


Fig.9. Box Plot of Closing Prices

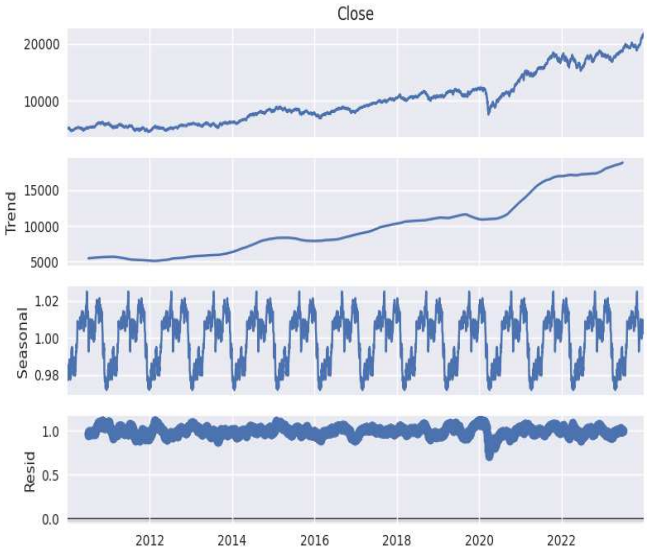


Fig. 10. Decomposition of Time Series - Close Price

Authorized licensed use limited to: ANNA UNIVERSITY. Downloaded on November 02,2026 at 04:09:53 UTC from IEEE Xplore. Restrictions apply.

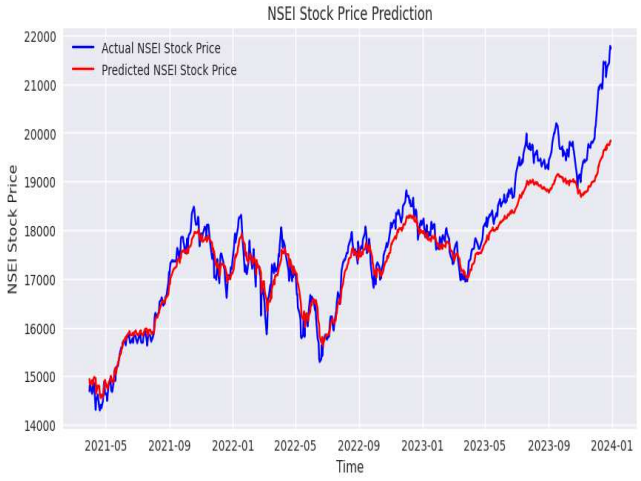


Fig. 12. NSEI Stock Price Prediction using LSTM

NSEI Stock Price Prediction using LSTM shown in Fig. 12. This plot compares the actual NSEI stock prices (in blue) with the predicted prices from the model (in red). The close alignment between the two lines suggests that the model has captured the trend and fluctuations in the stock prices reasonably well, although there are some deviations, especially in more volatile periods.

V. Conclusion

This paper demonstrated the effectiveness of XGBoost, an advanced gradient boosting algorithm, in predicting stock market volatility using historical time series data from the NSEI index, and compared its performance against the ARIMA model. XGBoost significantly outperformed ARIMA, achieving an MSE of 0.004, RMSE of 0.015, and an R^2 score of 0.98, showcasing excellent predictive accuracy and adaptability to market dynamics. In contrast, ARIMA, while capable of modeling linear relationships, was limited in handling the complex, non-linear patterns of financial data, resulting in an MSE of 0.72, RMSE of 0.41, and an R^2 score of 0.80. Additionally, a benchmark LSTM model achieved an MSE of 0.52, RMSE of 0.22, and an R^2 score of 0.45, indicating moderate performance in this context. XGBoost leveraged key features such as 'Open', 'High', 'Low', 'Close', 'Adj Close', and 'Volume,' with feature importance analysis validating their significance in the prediction process. These results highlight the robustness and flexibility of machine learning models like XGBoost in capturing financial volatility, surpassing traditional methods like ARIMA. Future research could explore hybrid models that combine the strengths of ARIMA and XGBoost or incorporate deep learning techniques and external economic indicators to further enhance predictive accuracy and applicability in financial forecasting.

References

[1]. R. Audrino, F. Sigrist, and D. Ballinari, "The impact of sentiment and attention measures on stock market volatility," in International Journal of Forecasting, vol. 36, no. 2, pp. 334–357, Apr. 2020, doi: 10.1016/j.ijforecast.2019.05.010.

[2]. N. D. Nagapoojitha, "Stock market predictions using machine learning techniques," in International Journal of Scientific Research in Engineering and Management, vol. 8, no. 7, pp. 1–13, Jul. 2024, doi: 10.55041/IJSREM36812.

[3]. H. Sun, "An example of machine learning-based multifactor dynamic quantitative stock picking models," in Dean & Francis, vol. 1, no. 7, Jun. 2024, doi: 10.61173/7n8qt956.

[4]. S. Modi and V. P. Upadhyay, "Optimizing stock market forecasts: The role of AI and hybrid models in predictive analytics, in International Journal of Automation, Artificial Intelligence and Machine Learning, vol. 4, no. 1, Jun. 2024, doi: 10.61797/ijaaiml.v4i1.329.

[5]. E. Idowu, "Advancements in financial market predictions using machine learning techniques, in Artificial Intelligence and Machine Learning, Jul. 2024, doi: 10.20944/preprints202407.1075.v1.

[6]. Sharma, T., Prasad, S.K. Enhancing cybersecurity in IoT networks: SLSTM-WCO algorithm for anomaly detection. Peer-to-Peer Netw. Appl. (2024). https://doi.org/10.1007/s12083-024-01712-z.

[7]. T. Sharma and S. K. Prasad, "Enhancing cybersecurity in IoT networks: SLSTM-WCO algorithm for anomaly detection," in Peer-to-Peer Networking and Applications, 2024, doi: 10.1007/s12083-024-01712-z.

[8]. J. Cui, Y. Zhang, and Y. Liu, "Machine learning for risk assessment in financial market forecasting," in Journal of Computing and Electronic Information Management, 2024, doi: 10.54097/0i9ppln6.

[9]. N. Agnihotri and S. K. Prasad, "Accuracy enhancement with artificial neural networks for bipolar disorder prediction," Indonesian Journal of Electrical Engineering and Computer Science, vol. 32, no. 3, pp. 1695–1702, Dec. 2023, doi: 10.11591/ijeecs.v32.i3.pp1695-1702.

[10]. D. David, W. Praveenraj, M. Pandey, and M. Victor, "Time series forecasting of stock market volatility using LSTM networks," in Proceedings of the IEEE International Conference on Computing Applications and Knowledge Management (ICCAKM), 2023, doi: 10.1109/iccakm58659.2023.10449596.

[11]. R. Zhu and J. Chen, "XGBoost and CNN-LSTM hybrid model with attention-based stock prediction," in Proceedings of the IEEE International Conference on Emerging Technologies and Computing Innovations (ICETCI), 2023, doi: 10.1109/ICETCI57876.2023.10176988.

[12]. H. Garg, R. Rana and S. K. Prasad, "A Model to Detect Face Mask Using Deep Learning," 2022 4th International Conference on Advances in Computing, Communication Control and Networking (ICAC3N), Greater Noida, India, 2022, pp. 944-949, doi: 10.1109/ICAC3N56670.2022.10073979.

[13]. D. Shah, H. Isah, and F. Zulkernine, "Stock market analysis: A review and taxonomy of prediction techniques," in International Journal of Financial Studies, vol. 7, no. 2, p. 26, 2019, doi: 10.3390/ijfs7020026.

[14]. K. Gopal, S. K. Prasad and A. Kumari, "Text and Facial Features Recognition using Machine Intelligence," Aug, 2022 International Conference on Machine Learning, Big Data, Cloud and Parallel Computing (COM-IT-CON), 2022, pp. 333-337, doi: 10.1109/COM-IT-CON54601.2022.9850445.

[15]. Y. Wang and Y. Guo, "Forecasting method of stock market volatility in time series data based on mixed model of ARIMA and XGBoost," in China Communications, vol. 17, no. 3, pp. 205-221, March 2020, doi: 10.23919/JCC.2020.03.017.

[16]. P. H. Vuong, T. T. Dat, T. K. Mai, P. H. Uyen, and P. The Bao, "Stock-Price Forecasting Based on XGBoost and LSTM," Comput. Syst. Sci. Eng., vol. 40, no. 1, pp. 237-246, 2022. https://doi.org/10.32604/csse.2022.0176855.

[17]. A. K. Dwivedi and S. K. Prasad, "Security enhancement scheduling model for IoT-based smart cities through machine learning method," *Advanced Control for Applications*, vol. 6, no. 4, Dec. 2024, doi: 10.1002/adc2.235.

[18]. P. Sharma and M. K. Jain, "Stock Market Trends Analysis using Extreme Gradient Boosting (XGBoost)," 2023 International Conference on Computing, Communication, and Intelligent Systems (ICCCIS), Greater Noida, India, 2023, pp. 317-322, doi: 10.1109/ICCCIS60361.2023.10425722.

Authorized licensed use limited to: ANNA UNIVERSITY. Downloaded on November 02,2026 at 04:09:53 UTC from IEEE Xplore. Restrictions apply.